

# Machine Learning for Omics Integration - Day 2 Notes

## Unsupervised Multi-Omics Integration

### Course Notes

December 17, 2024

## Contents

0.1	Key Challenges with Binary Data . . . . .	1
0.2	Unsupervised Machine Learning Philosophy . . . . .	2
0.3	MOFA: Multi-Omics Factor Analysis . . . . .	2
0.4	MOFA vs MOFA+: Detailed Comparison . . . . .	3
0.5	Model Evaluation . . . . .	6
<b>1</b>	<b>Artificial Neural Networks for Multi-Omics Integration</b>	<b>8</b>
1.1	Challenges of Deep Learning in Life Sciences . . . . .	8
1.2	General Principles of Artificial Neural Networks (ANN) . . . . .	8
1.3	Deep Learning Architectures for Multi-Omics . . . . .	9
1.4	Training Neural Networks . . . . .	9
1.5	Multi-Omics Deep Learning Applications . . . . .	10
1.6	Limitations and Alternatives . . . . .	10
1.7	Gradient Descent and Backpropagation . . . . .	10
<b>2</b>	<b>ANN from Scratch: Problem Formulation</b>	<b>12</b>
2.1	Key Insight: Neural Networks are About Derivatives, Not Black Boxes . . . . .	12
2.2	Autoencoders vs MOFA: Comparative Analysis . . . . .	15

*Continuation of Day 1 laboratory session (lab 2)*

### 0.1 Key Challenges with Binary Data

#### 0.1.1 Problem with Binary/Sparse Data

- **Binary data** (e.g., mutation data with mostly 0s and 1s) presents significant challenges
- **DIABLO** has known difficulties handling such sparse binary datasets
- **Recommendation:** Check mixOmics discussion forum for best practices and workarounds

*Day 2 lecture ## Pre-analysis Strategy: MOFA*

#### 0.1.2 Why Run MOFA First?

- **Purpose:** Understanding relationships between different omics layers before integration
- **Benefit:** Provides insights into data structure and inter-omics correlations
- **Strategy:** Use MOFA as exploratory analysis prior to supervised methods

## 0.2 Unsupervised Machine Learning Philosophy

### 0.2.1 The “Fishing Expedition” Approach

- **Concept:** Unsupervised ML is like a “fishing expedition”
- **No prior hypothesis:** We don’t understand the biological hypothesis beforehand
- **Discovery-driven:** Let the data reveal patterns and relationships

**Key Reference:** “*A hypothesis is a liability*” - article published in Genome Biology  
[Link to be added]

## 0.3 MOFA: Multi-Omics Factor Analysis

### 0.3.1 Overview

- **MOFA & MOFA+:** Leading examples of unsupervised multi-omics integration
- **Methodological approach:** Hybrid of PLS/CCA and Bayesian methods
- **Applications:** Widely used for discovering latent factors across omics layers

### 0.3.2 Core Concept: Factor Analysis

#### 0.3.2.1 What is Factor Analysis?

- **Central idea:** All observed data (gene expression, methylation, mutations) are generated by a few **latent variables** (factors/vectors)
- **Goal:** Learn these hidden factors from observed data
- **Process:** Start from observed data → infer hidden factors that explain the patterns

**0.3.2.2 Mathematical Foundation: Matrix Factorization** **Concept:** Decomposing the original data matrix into multiple component matrices

$$X_{ij} = U_{ik} \times V_{kj}$$

Where: - **X**: Original data matrix (samples × features) - **U**: Factor loadings matrix (samples × factors)  
- **V**: Factor weights matrix (factors × features) - **k**: Number of latent factors

### 0.3.3 MOFA Mathematical Framework

**0.3.3.1 Multi-View Factor Model** For multiple omics datasets, MOFA extends the basic factor model:

$$X_{ij}^{(m)} = \sum_{k=1}^K Z_{ik} \cdot W_{kj}^{(m)} + \epsilon_{ij}^{(m)}$$

**Where:** -  $\hat{\mathbf{X}}^{(m)}$ : Data matrix for omics type m (samples × features) - **Z**: Shared latent factor matrix (samples × factors) -  $\hat{\mathbf{W}}^{(m)}$ : Factor loadings for omics m (factors × features) - **K**: Number of latent factors -  $\hat{\epsilon}^{(m)}$ : Noise term for omics m

**0.3.3.2 Bayesian Formulation** MOFA uses a **Bayesian approach** with prior distributions:

**Factor Prior:**

$$Z_{ik} \sim \mathcal{N}(0, 1)$$

**Loading Prior (with sparsity):**

$$W_{kj}^{(m)} \sim \mathcal{N}(0, (\alpha_k^{(m)})^{-1})$$

Precision Prior (Automatic Relevance Determination):

$$\alpha_k^{(m)} \sim \text{Gamma}(a_0, b_0)$$

**0.3.3.3 Likelihood Functions** For continuous data (e.g., gene expression):

$$X_{ij}^{(m)} | Z, W^{(m)} \sim \mathcal{N} \left( \sum_{k=1}^K Z_{ik} W_{kj}^{(m)}, (\tau^{(m)})^{-1} \right)$$

For count data (e.g., RNA-seq):  $\text{X}(m)_{ij} | Z, W(m) \sim \text{Poisson}(\exp(\Sigma Z_{ik} \times W(m)_{kj}))$

For binary data (e.g., mutations):  $\text{X}(m)_{ij} | Z, W(m) \sim \text{Bernoulli}((\Sigma Z_{ik} \times W(m)_{kj}))$

Where is the sigmoid function:  $(x) = 1/(1 + e^{-x})$

**0.3.3.4 Variational Inference** MOFA uses **variational Bayes** to approximate the posterior distribution:

**Objective Function (ELBO):**

$$\mathcal{L} = \mathbb{E}_q[\log p(X, Z, W, \alpha)] - \mathbb{E}_q[\log q(Z, W, \alpha)]$$

Where: - **p(X, Z, W, alpha)**: Joint probability of data and parameters - **q(Z, W, alpha)**: Variational approximation to posterior - **ELBO**: Evidence Lower BOund (maximized during training)

### 0.3.4 Key Advantages of MOFA

1. **Missing Value Compensation:** Values missing in one omics layer can be compensated by information from other omics layers
2. **Variance Explanation:** Uses R<sup>2</sup> to quantify how much variance is explained by the model
3. **Cross-omics Discovery:** Identifies shared and unique factors across different data types

### 0.3.5 Applications

#### 0.3.5.1 scNMT Study

- **Example application:** Single-cell multi-omics integration
- **Reference:** scNMT paper [Link to be added]
- **Demonstrates:** Practical utility in real biological datasets

## 0.4 MOFA vs MOFA+: Detailed Comparison

### 0.4.1 MOFA (Multi-Omics Factor Analysis)

#### 0.4.1.1 Core Methodology

- **Statistical Framework:** Bayesian factor analysis with group sparsity
- **Key Innovation:** Handles multiple omics datasets simultaneously
- **Mathematical Foundation:**
  - Assumes data is generated from a low-dimensional latent space
  - Uses variational inference for model fitting
  - Incorporates automatic relevance determination (ARD) for factor selection

#### 0.4.1.2 MOFA Architecture

```
Input: Multiple omics matrices (RNA-seq, ATAC-seq, Methylation, etc.)  
↓  
Latent Factor Model: Z (samples × factors)  
↓  
Factor Loadings: W_m (factors × features) for each omics m  
↓  
Reconstruction: X_m = Z × W_m + noise
```

#### 0.4.1.3 Key Features of MOFA

1. **Factor Interpretability:** Each factor can be interpreted biologically
2. **Sparsity:** Automatically selects relevant features and factors
3. **Uncertainty Quantification:** Provides confidence intervals for estimates
4. **Missing Data Handling:** Naturally accommodates missing observations

### 0.4.2 MOFA+ (MOFA Plus)

#### 0.4.2.1 Major Improvements Over MOFA

- **Scalability:** Handles much larger datasets (>10,000 samples)
- **GPU Acceleration:** Faster computation using GPU implementations
- **Enhanced Flexibility:** Better handling of different data types and structures
- **Improved Convergence:** More robust optimization algorithms

#### 0.4.2.2 New Features in MOFA+

1. **Stochastic Variational Inference:** Enables mini-batch processing
2. **Non-Gaussian Likelihoods:** Better modeling of count data, binary data
3. **Smoothness Constraints:** For spatial/temporal data integration
4. **Transfer Learning:** Pre-trained models can be applied to new datasets

### 0.4.3 When to Use MOFA vs MOFA+

Aspect	MOFA	MOFA+
Dataset Size	< 5,000 samples	> 5,000 samples
Data Types	Continuous/Gaussian	Mixed data types
Computational Resources	CPU-friendly	Requires GPU for large data
Interpretability	High (simpler model)	High (with more complexity)
Development Status	Mature, stable	Active development

### 0.4.4 MOFA/MOFA+ Workflow

**0.4.4.1 1. Data Preprocessing** *Key preprocessing steps include:* - Log-transformation for count data - Feature filtering (highly variable features) - Normalization across samples - Quality control checks

**0.4.4.2 2. Model Training** *Basic MOFA model workflow involves:* - Creating MOFA object from data list - Setting model options (number of factors, etc.) - Configuring training parameters - Running the model training process

**0.4.4.3 3. Model Analysis**

- **Factor inspection:** Which factors explain most variance?
- **Feature loadings:** Which genes/features drive each factor?

- **Sample scores:** How do samples project onto factors?
- **Variance decomposition:** How much variance per omics layer?

#### 0.4.5 Biological Interpretation of MOFA Results

##### 0.4.5.1 Factor Types

1. **Shared Factors:** Active across multiple omics layers
  - Often represent fundamental biological processes
  - Examples: Cell cycle, differentiation states, stress responses
2. **Specific Factors:** Active in single omics layer
  - Capture omics-specific technical or biological variation
  - Examples: RNA processing effects, chromatin accessibility patterns

##### 0.4.5.2 Downstream Analysis

- **Pathway Enrichment:** Gene set analysis on factor loadings
- **Cell Type Identification:** Factor scores as features for clustering
- **Temporal Analysis:** Factor dynamics across time points
- **Clinical Association:** Correlate factors with phenotypes

#### 0.4.6 Advantages of MOFA Approach

##### 0.4.6.1 Over Traditional Methods

1. **vs PCA:** Handles multiple data types simultaneously
2. **vs CCA:** No requirement for paired samples across all omics
3. **vs Concatenation:** Accounts for different scales and noise levels
4. **vs Individual Analysis:** Identifies shared regulatory mechanisms

##### 0.4.6.2 Statistical Benefits

- **Dimensionality Reduction:** From thousands of features to ~10-50 factors
- **Noise Reduction:** Separates signal from technical noise
- **Integration:** Leverages complementary information across omics
- **Flexibility:** Accommodates different experimental designs

#### 0.4.7 Limitations and Considerations

##### 0.4.7.1 MOFA Limitations

1. **Linear Assumptions:** May miss non-linear relationships
2. **Factor Interpretation:** Requires biological expertise
3. **Hyperparameter Tuning:** Number of factors needs careful selection
4. **Computational Complexity:** Can be slow for very large datasets

##### 0.4.7.2 Best Practices

- **Factor Number Selection:** Use model selection criteria (ELBO, cross-validation)
- **Feature Selection:** Pre-filter for highly variable features
- **Batch Effects:** Address technical confounders before analysis
- **Validation:** Replicate findings in independent cohorts

#### 0.4.8 Case Study Applications

##### 0.4.8.1 Single-Cell Multi-Omics (scNMT-seq)

- **Data:** Single-cell RNA, DNA methylation, chromatin accessibility

- **Findings:** Identified developmental trajectories and regulatory relationships
- **Impact:** Revealed cell-type-specific regulatory mechanisms

#### 0.4.8.2 Cancer Multi-Omics

- **Data:** Gene expression, copy number, methylation, mutation
- **Findings:** Discovered pan-cancer molecular subtypes
- **Clinical Relevance:** Biomarkers for treatment stratification

#### 0.4.8.3 Population Studies

- **Data:** Multi-omics across large population cohorts
- **Findings:** Environmental and genetic factors affecting molecular profiles
- **Applications:** Precision medicine and risk prediction

### 0.5 Model Evaluation

#### 0.5.1 R<sup>2</sup> (R-squared)

- **Definition:** Proportion of variance in the data explained by the model
- **Range:** 0-1 (or 0-100%)
- **Interpretation:** Higher R<sup>2</sup> indicates better model fit and factor explanatory power

**0.5.1.1 Mathematical Formulation of R<sup>2</sup>** For each omics layer m and factor k, the variance explained is:

$$R_{mk}^2 = \frac{\text{Var}(\text{Predicted}_{mk})}{\text{Var}(\text{Observed}_m)}$$

Total variance explained across all factors for omics m:

$$R_m^2 = \sum_{k=1}^K R_{mk}^2$$

Overall model R<sup>2</sup> (across all omics):

$$R_{\text{total}}^2 = \frac{1}{M} \sum_{m=1}^M R_m^2$$

#### 0.5.2 MOFA-Specific Evaluation Metrics

**0.5.2.1 Variance Decomposition** The total variance in the data can be decomposed as:

$$\text{Var}(X^{(m)}) = \text{Var}_{\text{explained}} + \text{Var}_{\text{noise}}$$

Where:

$$\text{Var}_{\text{explained}} = \text{Var} \left( \sum_{k=1}^K Z_{ik} W_{kj}^{(m)} \right)$$

$$\text{Var}_{\text{noise}} = \text{Var}(\epsilon_{ij}^{(m)})$$

**0.5.2.2 Factor-Specific Variance Contribution** For factor k in omics m:

$$\text{Contribution}_{mk} = \frac{\text{Var}(Z_{ik}W_{kj}^{(m)})}{\text{Var}(X_{ij}^{(m)})}$$

**0.5.2.3 Evidence Lower Bound (ELBO)** The ELBO objective function can be decomposed as:

$$\mathcal{L} = \underbrace{\mathbb{E}_q[\log p(X|Z, W)]}_{\text{Reconstruction}} - \underbrace{D_{KL}(q(Z)||p(Z))}_{\text{Factor Regularization}} - \underbrace{D_{KL}(q(W)||p(W))}_{\text{Loading Regularization}}$$

**Where:** - **Reconstruction term:** How well the model reconstructs the original data - **KL divergence terms:** Regularization preventing overfitting - **D\_KL:** Kullback-Leibler divergence measuring difference between distributions

### 0.5.3 Model Selection Criteria

**0.5.3.1 ELBO-Based Model Selection** Compare models with different numbers of factors K:

$$\text{Best } K = \arg \max_K \mathcal{L}(K)$$

**0.5.3.2 Cross-Validation for Factor Selection** **K-fold CV procedure:** 1. Split data into K folds 2. For each fold i: Train MOFA on remaining folds, test on fold i 3. Compute CV error:

$$\text{CV Error} = \frac{1}{K} \sum_{i=1}^K \|\hat{X}^{(i)} - X^{(i)}\|^2$$

4. Select number of factors that minimize CV error

**0.5.3.3 Factor Activity Measure** Sparsity of factor k in omics m:

$$\text{Activity}_{mk} = \frac{\text{Number of non-zero } W_{kj}^{(m)}}{\text{Total number of features in omics m}}$$

### 0.5.4 Biological Validation Metrics

**0.5.4.1 Gene Set Enrichment** For factor k, compute enrichment p-value:

$$p_{\text{enrichment}} = P(\text{overlap} \geq \text{observed} | \text{random})$$

Using hypergeometric distribution:

$$P(X = x) = \frac{\binom{K}{x} \binom{N-K}{n-x}}{\binom{N}{n}}$$

Where: - **N:** Total genes in background - **K:** Genes in pathway  
- **n:** Genes associated with factor - **x:** Overlap between factor and pathway

**0.5.4.2 Factor Reproducibility** Correlation between factors across independent datasets:

$$\text{Reproducibility} = \text{cor}(Z_{\text{dataset1}}, Z_{\text{dataset2}})$$

Good reproducibility: correlation > 0.7

Sidenote: read MEFISTO method paper published in nature methods

# 1 Artificial Neural Networks for Multi-Omics Integration

## 1.1 Challenges of Deep Learning in Life Sciences

### 1.1.1 Major Limitations

- **Difficult application to real life science projects:** Most biological data (NGS, tabular omics data) doesn't naturally fit deep learning architectures
- **Lack of sufficient data:** Life sciences typically have small sample sizes compared to image/text domains
- **Simpler methods often outperform:** Traditional ML methods (LASSO, Random Forest) frequently achieve better results
- **Interpretability issues:** Black-box nature conflicts with biological understanding needs

### 1.1.2 Exceptions Where DL Excels

- **Single-cell omics:** Large cell numbers provide sufficient data
- **Medical imaging:** Microscopy, radiology, pathology images
- **Sequence analysis:** Protein/DNA sequence prediction tasks

## 1.2 General Principles of Artificial Neural Networks (ANN)

### 1.2.1 Mathematical Foundation

Artificial Neural Networks are **mathematical functions**  $Y = f(X)$  with a special architecture characterized by:

$$f(X) = f^{(L)}(f^{(L-1)}(\dots f^{(1)}(X)))$$

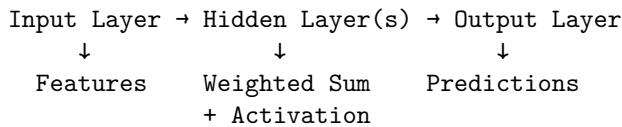
Where each layer  $l$  computes:

$$h^{(l)} = g(W^{(l)} \cdot h^{(l-1)} + b^{(l)})$$

**Components:** -  $W^{(l)}$ : Weight matrix for layer  $l$  -  $b^{(l)}$ : Bias vector for layer  $l$   
-  $g(\cdot)$ : Activation function (non-linear) -  $h^{(l)}$ : Hidden layer activations

### 1.2.2 Network Architecture

#### 1.2.2.1 Basic Structure



#### 1.2.2.2 Layer-by-Layer Process

1. **Input Layer:** Raw features (gene expression, methylation, etc.)
2. **Hidden Layer(s):**
  - Weighted sum:  $z = \sum_i w_i x_i + b$
  - Activation function:  $a = g(z)$
3. **Output Layer:** Final predictions or classifications

### 1.2.3 Activation Functions

#### 1.2.3.1 Common Activation Functions

Function	Formula	Range	Use Case
Sigmoid	$(x) = 1/(1+e^{-x})$	(0, 1)	Binary classification output
ReLU	$\text{ReLU}(x) = \max(0, x)$	$[0, \infty)$	Hidden layers (most common)
Tanh	$\tanh(x) = (e^x - e^{-x})/(e^x + e^{-x})$	(-1, 1)	Hidden layers (zero-centered)
Softmax	$\text{softmax}(x_i) = e^{x_i} / \sum e^{x_j}$	(0, 1)	Multi-class output

**1.2.3.2 Universal Approximation Theorem** Multi-layer neural networks with sufficient neurons can approximate any continuous function to arbitrary accuracy.

## 1.3 Deep Learning Architectures for Multi-Omics

### 1.3.1 1. Autoencoders

#### 1.3.1.1 Architecture

Input → Encoder → Latent Space → Decoder → Reconstruction  
 $(X) \quad \downarrow \quad (Z) \quad \downarrow \quad (X)$

**1.3.1.2 Mathematical Formulation** **Encoder:**  $Z = f_{\text{enc}}(X; \theta_{\text{enc}})$  **Decoder:**  $\hat{X} = f_{\text{dec}}(Z; \theta_{\text{dec}})$  **Loss Function:**  $L = \|X - \hat{X}\|^2 + \lambda R(\theta)$

#### 1.3.1.3 Multi-Omics Application

- **Input:** Concatenated omics data [RNA-seq | Methylation | Proteomics]
- **Latent space:** Compressed representation capturing shared patterns
- **Applications:** Dimensionality reduction, missing data imputation, batch correction

### 1.3.2 2. Multi-Modal Deep Learning

#### 1.3.2.1 Early Fusion (Concatenation)

RNA-seq → \  
Methylation → } → [Concatenated Vector] → Deep Network → Output  
Proteomics → /

#### 1.3.2.2 Late Fusion (Decision-Level)

RNA-seq → Network → \  
Methylation → Network → } → Fusion Layer → Output  
Proteomics → Network → /

## 1.4 Training Neural Networks

### 1.4.1 Backpropagation Algorithm

#### 1.4.1.1 Forward Pass

For each layer  $l$ :

$$z^{(l)} = W^{(l)}a^{(l-1)} + b^{(l)}$$

$$a^{(l)} = g(z^{(l)})$$

**1.4.1.2 Backward Pass (Gradient Descent for Error Minimization)** Loss gradient:  $\frac{\partial L}{\partial W^{(l)}} = \frac{\partial L}{\partial a^{(l)}} \frac{\partial a^{(l)}}{\partial z^{(l)}} \frac{\partial z^{(l)}}{\partial W^{(l)}}$

**Weight update:**  $W^{(l)} \leftarrow W^{(l)} - \eta \frac{\partial L}{\partial W^{(l)}}$

## 1.4.2 Loss Functions for Omics

**1.4.2.1 Regression Tasks** Mean Squared Error:  $L_{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$

**1.4.2.2 Classification Tasks** Cross-Entropy:  $L_{CE} = - \sum_{i=1}^n y_i \log(\hat{y}_i)$

# 1.5 Multi-Omics Deep Learning Applications

## 1.5.1 1. Cancer Subtyping

- **Input:** Gene expression + Methylation + Copy number + Mutation data
- **Architecture:** Multi-modal autoencoder → Clustering in latent space
- **Output:** Cancer subtypes with clinical relevance

## 1.5.2 2. Drug Response Prediction

- **Input:** Cell line omics profiles + Drug molecular features
- **Architecture:** Deep neural networks with multi-omics fusion
- **Output:** IC50 values or binary response predictions

## 1.5.3 3. Single-Cell Multi-Omics Integration

- **Input:** scRNA-seq + scATAC-seq + scProteomics
- **Architecture:** Variational autoencoders with modality-specific encoders
- **Output:** Integrated cell embeddings for trajectory analysis

# 1.6 Limitations and Alternatives

## 1.6.1 When NOT to Use Deep Learning

- **Small datasets:** < 1000 samples typically insufficient
- **High interpretability requirements:** Use linear models instead
- **Simple relationships:** Traditional ML often sufficient and faster

## 1.6.2 Better Alternatives for Small Data

- **MOFA/MOFA+:** Factor analysis approaches
- **Kernel methods:** SVM with multiple kernel learning
- **Ensemble methods:** Random Forest, XGBoost
- **Network-based methods:** Graph-based integration approaches

# 1.7 Gradient Descent and Backpropagation

## 1.7.1 Mathematical Formulation of Gradient Descent

Gradient descent is the cornerstone optimization algorithm for training neural networks. The objective is to minimize the loss function  $L(\theta)$  by iteratively updating parameters  $\theta$  in the direction of steepest descent.

**Basic Gradient Descent Update Rule:**

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} L(\theta_t)$$

Where: -  $\theta_t$  are the parameters at iteration  $t$  -  $\eta$  is the learning rate (crucial hyperparameter) -  $\nabla_{\theta} L(\theta_t)$  is the gradient of the loss function with respect to parameters

### 1.7.2 Backpropagation Algorithm

Backpropagation computes gradients efficiently using the chain rule of calculus. For a neural network with layers  $l = 1, 2, \dots, L$ :

**Forward Pass:**

$$\begin{aligned} z^{(l)} &= W^{(l)} h^{(l-1)} + b^{(l)} \\ h^{(l)} &= g(z^{(l)}) \end{aligned}$$

**Backward Pass:**

$$\delta^{(l)} = \frac{\partial L}{\partial z^{(l)}}$$

For the output layer ( $l = L$ ):

$$\delta^{(L)} = \frac{\partial L}{\partial h^{(L)}} \odot g'(z^{(L)})$$

For hidden layers ( $l = L - 1, L - 2, \dots, 1$ ):

$$\delta^{(l)} = [(W^{(l+1)})^T \delta^{(l+1)}] \odot g'(z^{(l)})$$

**Parameter Gradients:**

$$\begin{aligned} \frac{\partial L}{\partial W^{(l)}} &= \delta^{(l)} (h^{(l-1)})^T \\ \frac{\partial L}{\partial b^{(l)}} &= \delta^{(l)} \end{aligned}$$

### 1.7.3 Numeric Implementation Considerations

**1. Learning Rate Selection:** - Too large: oscillations, divergence - Too small: slow convergence - Common strategies: learning rate schedules, adaptive methods (Adam, RMSprop)

**2. Gradient Descent Variants:**

**Stochastic Gradient Descent (SGD):**

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} L_i(\theta_t)$$

**Mini-batch Gradient Descent:**

$$\theta_{t+1} = \theta_t - \eta \frac{1}{|B|} \sum_{i \in B} \nabla_{\theta} L_i(\theta_t)$$

**Adam Optimizer:**

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \\ \theta_{t+1} &= \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \end{aligned}$$

#### 1.7.4 Learning Rate as Critical Hyperparameter

The learning rate  $\eta$  significantly impacts training dynamics:

**Adaptive Learning Rate Strategies:** - **Step decay:**  $\eta_t = \eta_0 \cdot \gamma^{\lfloor t/s \rfloor}$  - **Exponential decay:**  $\eta_t = \eta_0 \cdot e^{-\lambda t}$   
 - **Cosine annealing:**  $\eta_t = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min})(1 + \cos(\frac{\pi t}{T}))$

**Learning Rate Scheduling for Multi-Omics:** Given the complexity and high dimensionality of omics data, careful learning rate management is crucial:  
 - Start with higher rates for initial feature learning  
 - Reduce rates as fine-tuning progresses  
 - Use validation loss plateaus to trigger rate reductions

#### 1.7.5 Practical Implementation for Multi-Omics Integration

**Challenge-Specific Considerations:**

1. **High Dimensionality:** Use gradient clipping to prevent exploding gradients
2. **Batch Normalization:** Helps stabilize training with different omics scales
3. **Regularization:** L1/L2 penalties added to loss function:

$$L_{total} = L_{data} + \lambda_1 \|W\|_1 + \lambda_2 \|W\|_2^2$$

**Example: Multi-Modal Loss Function**

$$L_{total} = L_{genomics} + L_{proteomics} + L_{metabolomics} + L_{integration}$$

Where each component can have different learning rates:

$$\begin{aligned}\theta_{genomics} &= \theta_{genomics} - \eta_g \nabla L_{genomics} \\ \theta_{proteomics} &= \theta_{proteomics} - \eta_p \nabla L_{proteomics}\end{aligned}$$

## 2 ANN from Scratch: Problem Formulation

### 2.1 Key Insight: Neural Networks are About Derivatives, Not Black Boxes

**Core Principle:** Neural networks are fundamentally about **gradients and derivatives** - they are interpretable mathematical functions, not mysterious black boxes.

#### 2.1.1 Mathematical Prerequisites for Understanding ANNs

##### 2.1.1.1 Essential Calculus Knowledge Time Investment for Learning Derivatives:

Background	Learning Timeline	Focus Areas
No Calculus	2-3 weeks intensive study	Chain rule, partial derivatives, gradients
Basic Calculus	3-5 days focused review	Multivariate calculus, chain rule applications
Strong Math Background	1-2 days	Review notation, vector calculus

##### 2.1.1.2 Core Derivative Concepts for ANNs 1. Chain Rule (Most Critical)

$$\frac{d}{dx}[f(g(x))] = f'(g(x)) \cdot g'(x)$$

**Example for Neural Networks:**

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial z} \cdot \frac{\partial z}{\partial w}$$

**2. Partial Derivatives** For function  $f(x, y)$ :

$$\frac{\partial f}{\partial x} = \text{derivative with respect to } x, \text{ treating } y \text{ as constant}$$

### 3. Vector Gradients

$$\nabla f = \left[ \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]$$

#### 2.1.2 Minimal ANN Implementation ( 10 Lines of Code)

Simple Neural Network from Scratch:

*This would be implemented using Python with numpy for basic operations, including:* - **Sigmoid activation function** and its derivative - **Network initialization** with random weights - **Forward pass** through hidden and output layers

- **Backward pass** computing gradients via chain rule - **Weight updates** using gradient descent - **Training on XOR problem** as a classic non-linear example

*Key insight: The entire neural network training process relies heavily on derivative calculations for the backward pass and weight updates.*

#### 2.1.3 Why Understanding Derivatives Matters for Omics

##### 2.1.3.1 1. Interpretability

- **Gradient analysis** reveals which genes/features most influence predictions
- **Sensitivity analysis** shows how changes in expression affect outcomes
- **Feature importance** derived from partial derivatives

##### 2.1.3.2 2. Debugging and Optimization

- **Vanishing gradients:** When derivatives  $\rightarrow 0$ , learning stops
- **Exploding gradients:** When derivatives  $\rightarrow \infty$ , training becomes unstable
- **Learning rate tuning:** Based on gradient magnitude

##### 2.1.3.3 3. Multi-Omics Specific Applications Gradient-Based Feature Selection:

$$\text{Importance}_i = \left| \frac{\partial L}{\partial x_i} \right|$$

Cross-Omics Gradient Analysis:

$$\frac{\partial L}{\partial \text{Gene}_i} \text{ vs } \frac{\partial L}{\partial \text{Protein}_j}$$

#### 2.1.4 Learning Path Recommendations

##### 2.1.4.1 Week 1-2: Calculus Fundamentals

- **Khan Academy:** Derivatives and Chain Rule
- **3Blue1Brown:** “Essence of Calculus” YouTube series
- **Focus:** Chain rule, partial derivatives, gradients

##### 2.1.4.2 Week 3: Applied to Neural Networks

- **Implement:** Simple ANN from scratch (above code)
- **Understand:** How backpropagation uses chain rule
- **Practice:** Calculate gradients by hand for 2-layer network

#### 2.1.4.3 Week 4: Multi-Omics Applications

- Gradient-based interpretation of trained models
- Feature importance analysis using derivatives
- Cross-omics gradient correlation analysis

#### 2.1.5 Quick Self-Assessment

Can you compute these derivatives?

1.  $d/dx[(wx + b)]$  where  $(x) = 1/(1+e^{-x})$
2.  $/w[(y - (wx + b))^2]$
3. Chain rule for:  $L/w$  in a 3-layer network

If yes: You're ready for ANN implementation If no: Spend 1-2 weeks on calculus review

#### 2.1.6 The “10 Lines” Philosophy

**Key Point:** Once you understand derivatives, neural networks become surprisingly simple:  
- **Forward pass:** Matrix multiplications + activation functions  
- **Backward pass:** Chain rule application  
- **Weight update:** Gradient descent

The complexity isn't in the math - it's in:  
- Architecture design - Hyperparameter tuning  
- Data preprocessing - Biological interpretation

10 Lines of Code for ANN:

```
import numpy as np

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(x):
    return x * (1 - x)

# Network parameters
np.random.seed(42)
weights_input_hidden = np.random.uniform(-1, 1, (2, 3)) # 2 inputs, 3 hidden neurons
weights_hidden_output = np.random.uniform(-1, 1, (3, 1)) # 3 hidden, 1 output

# Training data (XOR problem)
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([[0], [1], [1], [0]])

# Training loop
for epoch in range(10000):
    # Forward pass
    hidden_layer = sigmoid(np.dot(X, weights_input_hidden))
    output_layer = sigmoid(np.dot(hidden_layer, weights_hidden_output))

    # Backward pass (the derivatives in action!)
    output_error = y - output_layer
    output_delta = output_error * sigmoid_derivative(output_layer)

    hidden_error = output_delta.dot(weights_hidden_output.T)
    hidden_delta = hidden_error * sigmoid_derivative(hidden_layer)

    # Update weights (gradient descent)
```

```

weights_hidden_output += hidden_layer.T.dot(output_delta) * 0.1
weights_input_hidden += X.T.dot(hidden_delta) * 0.1

print("Final predictions:", output_layer.round(3))

```

## 2.2 Autoencoders vs MOFA: Comparative Analysis

### 2.2.1 Autoencoder Fundamentals

**Definition:** Autoencoders are **unsupervised neural networks** designed to learn compressed representations of input data.

**Core Objective:** Make the encoder and decoder reconstruct the input as perfectly as possible through a compressed bottleneck layer.

#### 2.2.1.1 Architecture Comparison

Autoencoder:	Input → Encoder → Bottleneck → Decoder → Output ( Input)
MOFA:	Multi-Omics → Factor Analysis → Latent Factors → Reconstruction

### 2.2.2 Mathematical Relationship to MOFA

**Autoencoder Loss Function:**

$$L_{AE} = \|X - \text{Decoder}(\text{Encoder}(X))\|^2$$

**MOFA Reconstruction:**

$$L_{MOFA} = \|X^{(m)} - ZW^{(m)}\|^2 + \text{regularization}$$

**Key Similarity:** Both methods aim to summarize high-dimensional data (thousands of genes/features) into a low-dimensional representation (2-50 factors/latent dimensions).

### 2.2.3 When Autoencoders Are Useful vs Not Useful

#### 2.2.3.1 Autoencoders Excel At:

1. **Large datasets:** > 10,000 samples (more data than typical genomics)
2. **Non-linear relationships:** Can capture complex patterns MOFA might miss
3. **Image-like data:** Spatial genomics, microscopy images
4. **Missing data imputation:** Learn to fill gaps in multi-omics datasets

#### 2.2.3.2 Autoencoders Limitations in Genomics:

1. **Small sample sizes:** Most genomics studies < 1,000 samples
2. **Interpretability:** Black-box nature vs MOFA's interpretable factors
3. **Overfitting risk:** Too many parameters for genomics datasets
4. **Better alternatives exist:** MOFA, PCA, factor analysis often superior

### 2.2.4 MOFA vs Autoencoder Trade-offs

Aspect	MOFA	Autoencoder
Sample size	Works with < 500 samples	Needs > 1,000 samples
Interpretability	High (factor meanings)	Low (black box)
Multi-omics handling	Native support	Requires concatenation
Missing data	Natural handling	Requires preprocessing

Aspect	MOFA	Autoencoder
Non-linear patterns	Limited (linear model)	Excellent
Computational cost	Moderate	High (GPU needed)
Biological insight	High	Limited

### 2.2.5 Practical Recommendation

**For typical genomics projects:** Use MOFA first - Provides interpretable biological insights - Handles small sample sizes well - Better suited for multi-omics integration - Established in genomics literature

**Consider autoencoders when:** - Dataset > 5,000 samples - Strong non-linear relationships suspected - Working with spatial/imaging data - Need missing data imputation

### 2.2.6 Critical Insight: Context Matters

**“Autoencoder dimension reduction: somewhat useful, but there are better techniques”**

This reflects the reality that **simpler methods often outperform complex ones** in life sciences, especially with limited data.

**Deep learning limitations in critical applications:** - **Life-death situations:** Medical diagnostics require interpretability - **Safety-critical domains:** Earthquake prediction, drug discovery - **Regulatory approval:** Black-box models difficult to validate

**The “right tool for the job” principle:** Match method complexity to: - Data size - Interpretability requirements

- Domain constraints - Validation needs